

Movie Recommender System

...

MovieLens Dataset

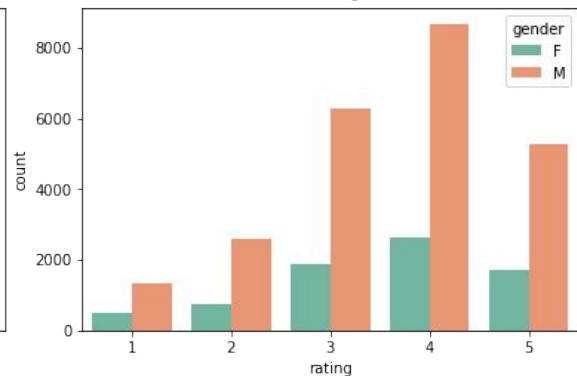
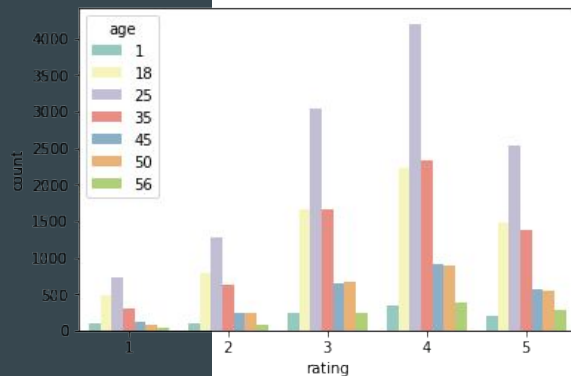



Yuyang Tian
Yuting Xu

Xiaoyun Wang
Don Park

Initial Data Analysis

Dataset

- 31620 ratings
- 1465 movies
- 2353 users
- Density: 0.917%



Collaborative Filtering

Reveal User v.s. Movie Interaction

Add Biases

Optimization

Latent Factor Model (SVD)

Baseline Predictor

Stochastic Gradient Descent

- Find the latent factors
 - Romance? Horror?
- Weigh the relationship between users and factors & movies and factors
- Derive unknown ratings by factors

Separate user and movie behavior

- User Bias
 - Rating scale of a user
 - Behavior of a user
- Movie Bias
 - Popularity of a movie
 - Selection bias

- Stochastic: Fast Convergence in practice
- Gradient Descent: Find optimal components to reconstruct a full matrix with the minimal loss against the original sparse matrix

Matrix Factorization

$$\hat{r}_{ui} = \mathbf{q}_i \mathbf{p}_u = \sum_f q_{if} p_{fu}$$

1. “Matrix factorization characterizes both movies and users by vector of factors inferred from item rating patterns.” [Koren]
2. How to fill the blank entries?

	user1	user2	user3	user4	≈		factor1	factor2	factor3	×	(-0.5)*(-2)+0.6*0.3+1.5*2.4=4.78(estimated)					
movie1	1		4			movie1	q11	q12	q13							
movie2		5				movie2	q21	q22	q23							
movie3				5		movie3	q31	q32	q33			factor1	p11	p12	-2	p14
movie4	4		2			movie4	q41	q42	q43			factor2	p21	p22	0.3	p24
movie5						movie5	q51	q52	q53			factor3	p31	p32	2.4	p34
movie6		3	4.78			movie6	-0.5	0.6	1.5			P = 3 rows × 4columns				
movie7				3		movie7	q71	q72	q73							
movie8	2					movie8	q81	q82	q83							
movie9		3				movie9	q91	q92	q93							
movie10			4		movie10	q10 1	q10 2	q10 3								
10 rows × 4 columns						Transpose Q = 10rows × 3columns										

Objective Function

$$\min_{q_u, p_i} \sum_{u,i} (r_{ui} - \mathbf{q}_i^T \mathbf{p}_u)^2 + \lambda (\|\mathbf{q}_i\|^2 + \|\mathbf{p}_u\|^2)$$

Cost function: min(SSE + Regularization)

SSE: how well the mode is fit

Regularization: avoid overfit

Our goal is to find the Q and P

The idea: minimum of the objective function

Overall: we used known rating to get P,Q then predict rating.

.

$$r_{xi} = \underbrace{\mu}_{\text{Overall mean rating}} + \underbrace{b_x}_{\text{Bias for user } x} + \underbrace{b_i}_{\text{Bias for movie } i} + \underbrace{q_i \cdot p_x}_{\text{User-Movie interaction}}$$

Stochastic gradient descent

- **Gradient descent optimization**

To minimize the objective function $f(x)$:

- find the direction in which decreases the fastest: negative gradient
- propose a new point $x' = x - \gamma \nabla_x f(x)$, where γ is the learning rate
- converges when every element of the gradient is zero (or, in practice, very close to zero).



- **Stochastic gradient descent**

—Minimize an objective function that has the form of a sum:

$$Q(\omega) = \frac{1}{n} \sum_{i=1}^n Q_i(\omega)$$
$$\omega = \omega - \gamma \nabla Q(\omega) = \omega - \gamma \frac{1}{n} \sum_{i=1}^n \nabla Q_i(\omega)$$

Problem:

large training sets are more computationally expensive.

—Approximated by a gradient at a single sample

$$\omega = \omega - \gamma \nabla Q_i(\omega)$$

Model Process

- Randomly select a sample r_{ui} from training set
- Predict \hat{r}_{ui} by $\hat{r}_{ui} = q_i^T p_u$
- Compute the associated prediction error $e_{ui} \stackrel{def}{=} r_{ui} - q_i^T p_u$
- Update the parameters by a learning rate in the opposite direction of the gradient

$$\begin{aligned} q_i &\leftarrow q_i + \gamma \cdot (e_{ui} \cdot p_u - \lambda \cdot q_i) \\ p_u &\leftarrow p_u + \gamma \cdot (e_{ui} \cdot q_i - \lambda \cdot p_u) \end{aligned}$$

- Repeat step 1-4, until find the optimal p_u, q_i
- Generate the prediction rating matrix

$$R_{pred} = Q^T \times P = [q_1, q_2, \dots, q_I]^T \times [p_1, p_2, \dots, p_U]$$

SVD Only

Configuration

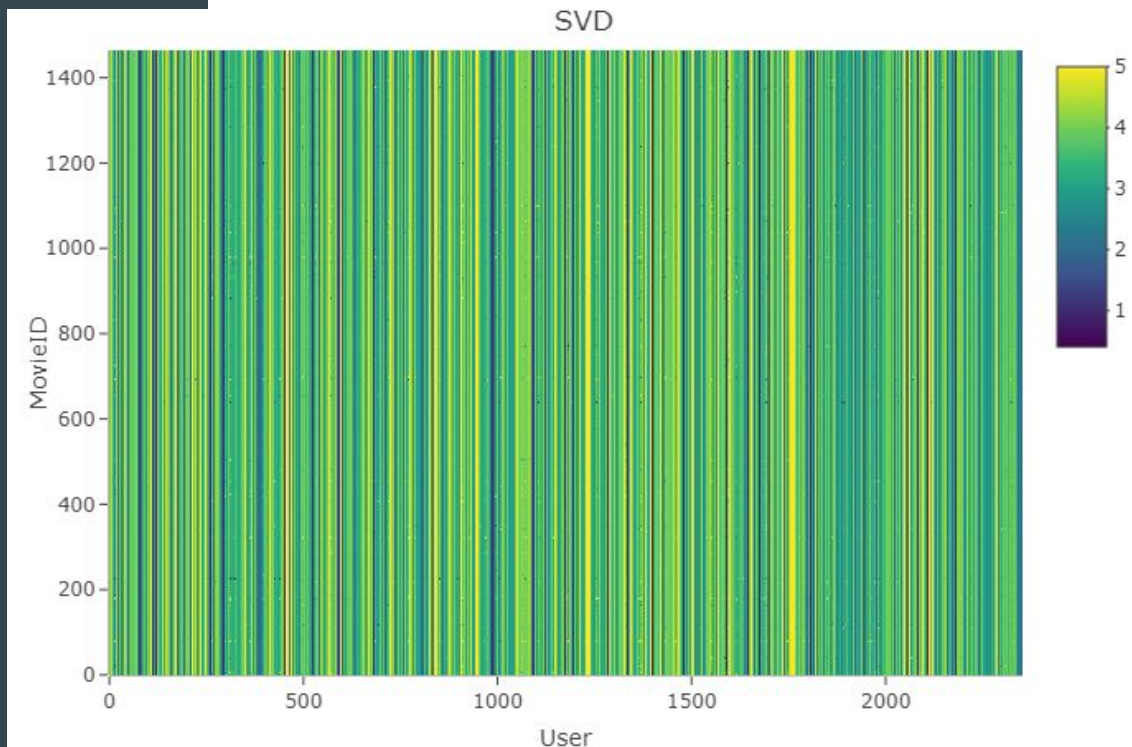
- # of latent factors: 100
- Fill unrated data with 0 (mean)
- Normalized then trained
- No calibration run (default configuration)

Result

RMSE = 1.429

Good start, but can do better.

Biased based on mean rating



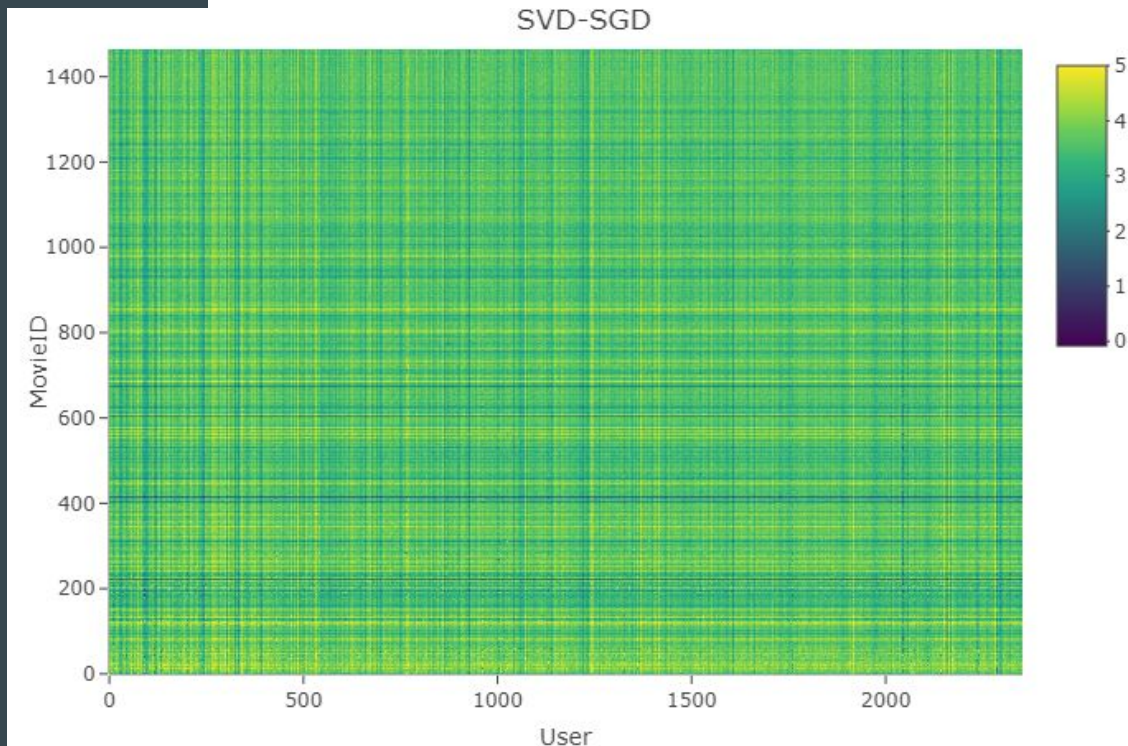
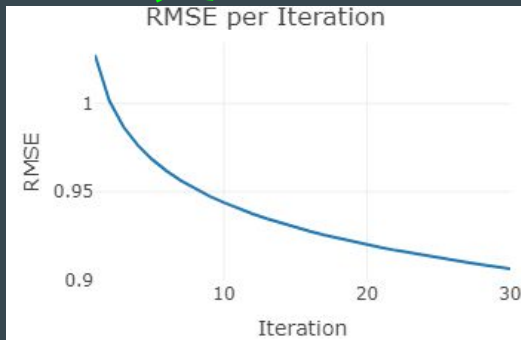
SVD with SGD

Configuration

- # of latent factors: 100
- # of epochs: 30
- Learning Rate: 0.005
- Regularization Term: 0.02
- Fill unrated data with 0 (mean)
- Normalized then trained

Result

RMSE = **0.927**



Summary

Summary

- Matrix Factorization can be used to predict user ratings
- SVD, an extension of Matrix Factorization, is a good model to predict user ratings
- SVD+SGD has improved performance over just SVD